

University of São Paulo  
Institute of Mathematics and Statistics  
Bachelor's in Computer Science

Breno Helfstein Moura

# Hashing Functions and Hash Tables

## A practical approach

São Paulo  
November de 2019

# Hashing Functions and Hash Tables

## A practical approach

Final undergraduate thesis for subject  
MAC0499 – Trabalho de Formatura Supervisionado.

Supervisor: Prof. Dr. Jose Coelho de Pina Junior  
[ Cosupervisor: Prof. Dr. Nina S. T. Hirata ]

São Paulo  
November de 2019

# Abstract

During this undergraduate thesis I will explain about two of the most fascinating and used ideas in Computer Science, hash functions and hash tables. I divided this thesis in three main parts:

- Hash functions
- Hash tables
- Applications

During the first part I explain why hash functions are an important idea in Computer Science, summarize some of the ideas Donald Knuth present on his book (The Art of Computer programming, Vol. 3) and use some metrics to evaluate what is a good hash function.

During the second part I talk about one of the most used data structures in computer programming, hash tables. I will explain what constitute a hash table, show some of the classic implementations of this data structure and explain some of the most used open addressing strategies. It is nice to observe here that although hash tables is a simple concept, there is still debates regarding this subject with no clear consensus on what is a state of the art hash table.

During the third, and last, part I will cite and explain some application of hash functions in computer science problems. I will explain Rabin-Karp, a string search algorithm that uses hashing and a solution to identify isomorphisms on trees using hashing functions.

I hope this is as fun to read for you as it was for me to write!

**Keywords:** hash functions, hash tables, collision-resolution, open-addressing.



# Contents


<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Hash Functions</b>	<b>3</b>
<b>3</b>	<b>Hash Tables</b>	<b>11</b>
3.1	Open Addressing . . . . .	11
3.1.1	Implementing a no collision hash table . . . . .	11
3.1.2	Formal definition . . . . .	14
3.1.3	Linear Probing . . . . .	14
3.1.4	Quadratic Probing . . . . .	14
3.1.5	When to resize an array . . . . .	14
3.1.6	How to delete an entry . . . . .	14
3.1.7	Practical advantages of Open Addressing . . . . .	14
3.2	Chaining Hashing . . . . .	14
3.2.1	Simple Chaining Hashing Algorithm . . . . .	14
3.2.2	Move to front . . . . .	16
3.2.3	When to resize an array . . . . .	16
3.2.4	How to delete an entry . . . . .	16
3.2.5	Practical advantages of Chaining Hashing . . . . .	16
<b>4</b>	<b>Applications</b>	<b>17</b>
4.1	3-sum problem . . . . .	17
<b>5</b>	<b>Conclusions</b>	<b>21</b>
<b>A</b>	<b>Apendix</b>	<b>23</b>



# Chapter 1

## Introduction

One of the most used data structures in computer science are dictionaries. Those need to suport the operations of inserting, fiding and deleting an element. If you think about it, this is one of the most executed tasks in many softwares. For example, when you have the list of numbers you last called on your cell phone and you want to know for each phone number, what is the person associated with it. A dictionary perfoms the task of inserting for each phone number the name of the person, and then you can retrieve that information finding, for a phone number who is the person associated with it.




**Figure 1.1:** *Example of a dictionary that associates phone numbers to contact names.*

Other use of a dictonary that you can think is to count the number of times you called a certain number. One of the most used implementation of dictionaries is with a hash table.

The implementation of a hash table always requires a hash function. This function usually takes the element that you want to hash, or key as it is usually called (In our example, the phone numbers), and “digest” it into a number. That number is then used to indentify the value (In our eaxmple, the contact names), in this structure that we call hash table.

An example of a hash function, that “digest” the phone numbers is the following:



**Figure 1.2:** Example of a hash function hash function that just take the last 2 digits of the phone number

As you can see this is a pretty simple function, it simply take the last 2 digits of each phone number. In this specific case, this is enough to uniquely identify each phone. We can imagine a function that can't uniquely identify each phone number, like getting just the last digit (in this case, Cauli's and Breno's numbers would have the same hash value), this will cause a collision in the table. Solving collisions in a hash table is a complete topic by itself, and it will be addressed in Chapter 3, Hash Tables.

Solving collisions is actually a very important topic in hash tables, and that is because the vast majority of hash functions will have collisions. To picture that we can remember the "Birthday Paradox", that is the conclusion that we only need 23 people in a room to have a chance greater than 50% of 2 or more people having the same birthday. In Donald Knuth's famous book, The Art of Computer Programming (Vol. 3, Chapter 6.4) [? ], he uses as an example a function from a 31-element set to a 41-element set, and from about  $10^{50}$  functions only about  $10^{43}$  give distinct values for each argument, that is about 1 in every 10 million functions. That shows that we will have collisions more often than not, so knowing how to deal with it is a major problem.

Hash functions and hash tables are among the most classic topics within computer science, yet is still one of the topics with most debate about what is state of the art. While the hash table was invented in 1953, widely discussed by Donald Knuth in his book, there are still many tweaks that can be made to boost its performance for specific use cases. One great example is F14, an open-source memory efficient hash table by Facebook <sup>1</sup>.

An example of lack of consensus in this area are the different hash functions and hash table implementations in different languages. There is no clear consensus on how to decide the size of a hash table, what are the tradeoffs of the collision-resolution algorithms or even what defines a good hash function. Hopefully, we got years of research on the topic to study and present a view on the subject, and that is what I am presenting throughout this undergraduate thesis.

<sup>1</sup>F14 is open sourced: <https://engineering.fb.com/developer-tools/f14/>




# Chapter 2

## Hash Functions

Outside computer science, the word “*hash*” in the english language means to “chop” or to “mix” something. This meaning is entirely related to what hash functions are supposed to do. hash functions are functions that are used to map data of an arbitrary size to data of a fixed size [? ].

They have wide applications in computer science, being used in information and data security, compilers, distributed systems and hardcore algorithms. During this chapter I first define and explain the basics of a hash function, then I give an intuition in some metrics of what is a good hash function, as discussed in the famous “*Red Dragon Book*” [? ] along with some reproduction of known results in the area.

The value extracted from the hash function for an object is usually called *Hash Value*. The hash value is usually, but not necessarily, smaller than the object that generated it. For example, we can have a hash function that takes Gigabytes or Terabytes files and return an 8 bytes hash value.



**Figure 2.1:** Example of a hash function from string to 4 bit integer.

To formalize a little, let's define a hash function as a function  $H$  that takes an element  $x \in X$  and has  $[0, M)$  as a codomain.

This is the same definition used by Donald Knuth [?] and some articles [?]. This definition makes sense for our case because we will be talking mostly about hash functions used in hash tables, and in that case we want integers that will be indexes in an array (as we will see later on). In other cases we may see hash functions value as strings, like for when we hash a string for password storage or when we use a hash function in files for check-sum (for when we are checking if two files are the same). For the goal of this thesis we will not be focusing on those functions, but it is important to notice that strings can also be abstracted to integers if we just look at the bytes.

For our specific case we are looking at a hash function that is good for the construction of hash tables, that is that is fast to calculate and minimizes the number of collisions. Depending on our goals we might want a different metric, for check-sums for example we may want a function that is very sensible to changes, and for passwords one that is very hard to find its inverse.

As said in Donald Knuth's book, we know that it is theoretically impossible to create a hash function that generates true random data from non random data in actual file, but we can do pretty close to that (or in some cases, even better). Donald Knuth describes 2 specific methods for simple hash function, named *division hashing* and *multiplicative hashing* techniques. As the name suggests, the first is based on division and the former on multiplication.

The division hashing method simply to represent the data as a number take the remainder of that number modulo a value. Supposing that we can represent the data as a non negative integer  $X$  the division hashing would be to choose a value  $M$  and the hashing function would be  $X \bmod M$ . The C++11 code would look as following:

```
1 unsigned_int divisionHashing(unsigned_int X, unsigned_int M) {
2     return X % M;
3 }
```

In general large prime numbers tend to be a good value to  $M$ , because if not we may have repetitions. One great example of this is if  $M$  is even, then the parity of hash value of  $X$  will match the parity of  $X$  (which will cause a bad distribution). The same pattern will happen in different intervals for different powers of 2.

For the multiplicative hashing, we can first imagine that the overflow is like a "natural" modulo operation (We also have methods to take the modulo without overflowing, to know more about that). Supposing that we can represent the data as a non negative integer  $X$ , the multiplication hashing would be to choose a value  $A$  that we multiply by  $X$  and then take the value modulo  $2^P$  (That is how is described in donald knuth book, by "*taking the leading bits of the least significant half of  $A * X$* "). The C++11 code would look as following:

```

1 unsigned_int multiplicativeHashing(unsigned_int X,
2                                   unsigned_int A,
3                                   unsigned_int P) {
4     return (A * X) << P;
5 }

```

In knuth’s book he restricts  $A$  to be relatively prime to  $w$ , being  $w$  the size of a “word” in the machine (which is  $MAX\_INT$  in our case). That definition is often useful if you can retrieve a value  $Y$  for a given hash value  $F$ . It is good to note here that if  $H(X) = F$  and  $H(Y) = F$ ,  $X$  is not necessarily equal to  $Y$ , as two keys can have the same hash value.

Here it is also good to note, we have many ways of converting non numerical data to non negative integers. In the end, it is all just a sequence bytes, that when readed in a specific way form another type of data, such as images or strings. For example, one way of transforming a string to a non negative integer is summing the ASCII value of its characters. The C++11 code for that would look as following:

```

1 unsigned_int convertStringToInteger(string str) {
2     unsigned_int hashValue = 0;
3     for (char c : str) {
4         hashValue += (int) c;
5     }
6     return hashValue;
7 }

```

We always use usigned integers for our non negative integer calculations due to the natural modulo operation of it on overflow cases. It is equivalent to having a  $\text{mod } 2^{32}$  every time it overflows (As we only look at the leading 32 bits). We can also use  $XOR$  function to mix numbers together.

There is also a very common type of hash functions that tend to work pretty well for strings [? ]. It is a “Superset” of multiplicative hash functions, or a generalization. The C++11 code would look as following:

```

1 unsigned_int hashForString(string str ,
2                             unsigned_int initialValue ,
3                             unsigned_int multiplier ,
4                             unsigned_int modulo) {
5     unsigned_int hash = initialValue;
6     for (char c : str) {
7         hash = (multiplier * hash + (int)c);
8     }
9     return hash % modulo;
10 }

```

The above function is very common for string hashing, and by just choosing a different initial value and multiplier we can have completely different hash functions. Although using summing or using XOR usually don't provide much difference, XOR is preferable due to the fact that we do not need to worry about overflow. Some values are of known hash functions, for example with *multiplier* = 33 and *initialValue* = 5381 generates *Bernstein hash djb2* [?] or *multiplier* = 31 and *initialValue* = 0 generates *Kernighan and Ritchie's hash* [?]. Those are famous functions and their values are not chosen randomly, as there are some factors that maximize the chance of producing a good hash function (remembering, good means low collision rate and fast computation). Those factors are:

- The multiplier should be bigger than the size of your alphabet (in our case usually 26 for english words or 256 for ASCII). That is the case because if it is smaller we can have wrong matches easier. For example, suppose that *multiplier* = 10 and *initialValue* = 0, we have  $H('ABA') = H('AAK') = 7225$  before taking the modulo operation.
- The multiplication by the multiplier should be easy to calculate with simple operations, such as bitwise operations and adding. That is quite intuitive as we want a hash function that is fast to calculate.
- The multiplier should be coprime with the modulo. That is because if not we will "cycle" hashes at a greater rate than the modulo (We can use some modular arithmetic to prove that). Usually prime numbers tend to be good multipliers.

Now that we know some good templates for producing hash functions let's try to find a concrete metric or formula that measures the quality of a hash function. Fortunately, the famous Red Dragon Book [?] has already proposed a formula to measure the quality of a hash function. The formula is the following:

$$\sum_{j=0}^{m-1} \frac{b_j(b_j + 1)/2}{(n/2m)(n + 2m - 1)}$$

Where  $n$  is the number of keys,  $m$  is the number of total slots and  $b_j$  is the number of keys in the  $j$ -th slot. The intuition for the numerator is the number of operations we will need to execute to find each element of the table. For example, we need 1 operation to find the first element, 2 to find the second, and so on. That means that we will end up with the following arithmetic progression. We know that a hash function that distributes the keys in a uniformly random distribution has expected bucket size of  $n/m$ , so we can calculate that the expected value of the numerator formula is  $(n/2m)(n + 2m - 1)$ . So that gives us a ratio of collisions (thinking just about operations to access a value) of "our" hash function with an "ideal" function. That means that a value close to 1.00 of the above formula is good, and values below to 1.00 means that we had less conflicts than an uniformly distributed random function.

From common data as used in Dragon Book and Strchr website [?] I will reproduce some tests with the previously cited hash functions.




Figure 2.2: Functions tested against a “small” table

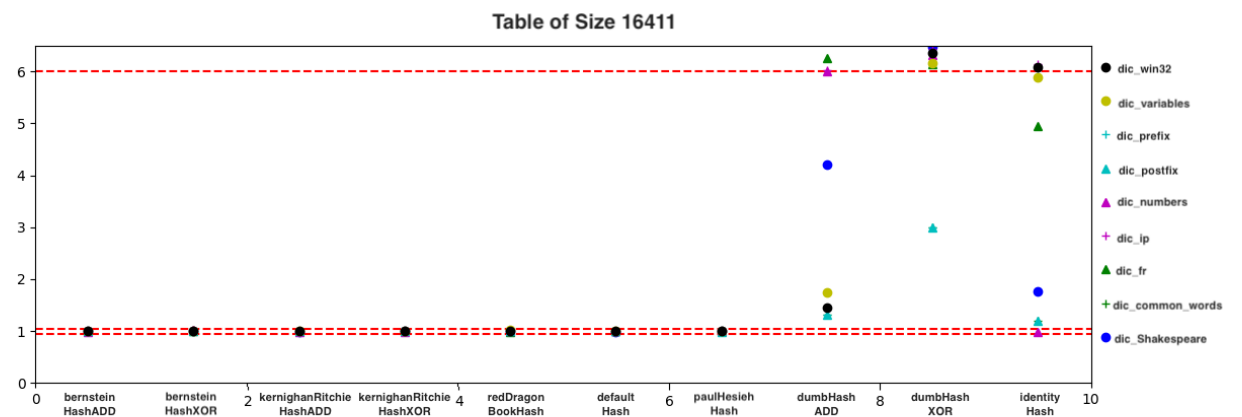


Figure 2.3: Functions tested against a “medium” sized table




Figure 2.4: Functions tested against a “large” table

The results are shown in the same way displayed in the Red Dragon Book, with Hash functions in the X axis and the ratio displayed in the y axis, with different identification for each file. I choose 3 sizes of tables to count collisions, a “small”, “medium” and a “large” table, the small table having a load factor (The percentage of the table occupied) of approximately  $\sim 0.5$ , the medium with  $\sim 0.05$  and large with  $\sim 0.005$ . It’s assumed that the modulo is not the responsability of the hash function, so all hash functions return values from 0 to  $2^{32} - 1$ , and the modulo is taken depending on the size of the table. From that we can already see that the load factor doesn’t make a good hash function bad, but expose problems of “bad” hash functions in some cases.

Another thing that it is important to notice from the graph is the red dotted lines. The top one is the “Upper” threshold, which results greater than 6 are just considered “Big”, as in some cases the ration exploded to values up to 200. The lower 2 red lines are in  $y = 1.05$  and  $y = 0.95$  which is the interval that we consider a hash function to have “Good” values.

The tests were made with 10 different hash functions, tested against 9 different files (Provided by strchr website). All of the code used to test this can be found in the github repo [? ]. The 10 hash functions are the following:

- **bernsteinHashADD**: The bernstein hash function described earlier. We use the given hash template adding the elements. The multiplier is 33 and initial value is 5381. In the end we XOR the bits of the hash with itself shifted 16 to the right (That is half of the bits with our implementation).
- **bernsteinHashXOR**: The same as above but substituting the first adding operator by the XOR operation.
- **kernighanRitchieHashADD** The Kernighan Ritchie Hash function described earlier. We use the given hash template adding the elements. The multiplier is 31 and initial value is 0.
- **kernighanRitchieHashXOR** The same as above but substituting the first adding operator by the XOR operation.
- **redDragonBookHash** The hash function tested in the red dragon book. It is described as x65599 in the book.
- **defaultHash** The default hash function of c++ standard template library.
- **paulHesiehHash** A fast hash function described by Paul Hesieh. It is fast to calculate and more complex than Knuth multiplicative or division Hashing.
- **dumbHashADD** A hash function that simply add all characters.
- **dumbHashXOR** A hash function that simply XOR all characters.

- **identityHash** A hash function that takes the first 4 bytes of the string.

We have a variety of hash functions, with all being considered “fast” hash functions. The files tested include common words in english and french, strings of some IP values, numbers, common variable names and words with common prefix and postfix.

First thing we can notice from those graphs is that changing the ADD function to XOR doesn’t make a good multiplicative hash function bad. Both are actually “Equivalent” given that we are also multiplying the values. For “dumbHashADD” and “dumbHashXOR” we can see clear differences, with “dumbHashXOR” being clearly worse. This can be explained by the cancelation property of XOR. We can see this example on the hash of this IP below:

$$\text{dumbHashXOR}("168.1.1.0") = \text{dumbHashXOR}("168.2.2.0") = \text{dumbHashXOR}("124.6.8.0")$$

We can see that many different IPs have the same hash value. More than that, XOR doesn’t increase the number of bits, so all the hashes will be of just 1 byte.

Other thing that we can notice is that “identity” hash is good or perfect in some cases. One obvious case that “identity” works perfectly is for numbers (We will have 0 collisions). Some languages, like python 3, use the identity function to calculate hash for integers, as it is very fast and produces no collisions. But we can see that for other cases, such as common prefix, it works terrible as we just get the first 4 bytes.

The most common multiplicative hash functions tend to work similarly well, being reasonably close to the ideal hash function in all cases.

As we can see, we don’t need a lot of complexity to make a good hash function for a hash table. We have some functions working better for some specific case, like identity function working well for numbers, but general functions already work well enough.

One thing that is important to say here is that hash function is a very vast topic, and here we just covered hash functions related to hash tables. Hash functions have applications in distributed systems (With consistent hashing), database indexing, caching, compilers (As it is explored in the red dragon book) and cryptography. Each application has different requirements and make some hash functions better than others.





# Chapter 3

## Hash Tables

Hash tables or hash maps is one of the most used applications of hash functions. It is actually so used in computer science that is almost impossible to talk about one without mentioning the other. This data structure consists in associating a *key* to a *value* in a table. That is, given a *key*, it can retrieve the correct *value* for it.

It is considered one of the possible (and one of the best) implementations of a dictionary. It usually has to implement the **INSERT**, **FIND** and **REMOVE** operations, that can be accessed from outside the dictionary. It usually implements a lot of other private methods.

This data structure is usually considered very useful among software engineers and computer scientists, although it usually has a linear worst case cost for retrieving, inserting and deleting a key-value pair. That is because it has a constant average cost for those operations.


Moreover, when talking about hash tables we have the problem of key collision, that is when two keys maps to the same hash value. As we saw in the previous chapter, this is collisions are more common than not. To solve that problem, we have several techniques that involves different tradeoffs. Those techniques are usually divided into two main categories, open addressing and separate chaining. Other problem to consider regarding this data structure is when to resize the hash table, to minimize the chance of collision and the use of memory. For this last one we usually consider a load factor,  $\alpha$ , that is the ratio of keys with the available slots.

It is also important to notice that hash tables have applications in different areas of computer science also, like compilers, caches and database indexing.

### 3.1 Open Addressing

#### 3.1.1 Implementing a no collision hash table

To start I will give an example of a hash table that has a perfect hash function, that has no collisions. For that example we will use open addressing, that basically means that all data will be contained in an array. The operations **INSERT**, **FIND** and **REMOVE**



**Figure 3.1:** Example of a hash table from string to string, more specifically name to number

would be very easy to implement. For the sake of simplicity, I will assume all the keys are strings. To start let's look at this simple class with dummy methods:

```

1 class HashTable {
2     vector< pair<string, int> > table;
3     int m, n;
4
5     HashTable() {
6         m = 16;
7         table.resize(m);
8         n = 0;
9     }
10
11     unsigned_int hashFunction(string s) {}
12
13     void insert(string key, int value) {}
14
15     int find(string key) {}
16
17     void remove(string key) {}
18
19 private:
20     void resizeIfNecessary() {}
21 }

```

As we can see it is pretty simple. The constructor builds a table of size 16, and for now we can assume a dynamic resizing every time we have 16 elements. Later on we will see that

this means that we resize every time the load factor,  $\alpha$ , is equal to 1.00. We also can note that at the table part we are storing a pair of key and value, not just value. This is because we may want to retrieve all pairs of the table (like in a regular dictionary). The pairs are usually unordered (If they are not ordered by chance ...), and the iterator has  $O(1)$  step. We will jump the implementation of `hashFunction`, as already saw plenty of it in the last chapter, so we will go right in for the implementation of **insert**:

```
1 void insert(string key, int value) {
2     unsigned int idx = hashFunction(key);
3     table[idx] = pair<string, int>(key, value);
4     n++;
5     resizeIfNecessary();
6 }
```

That is pretty simple, that is mostly because we will assume that we will never have a collision, so we just put the key on the position returned by the hash function. The method **find** is implemented as following:

```
1 int find(string key) {
2     unsigned int idx = hashFunction(key);
3     if (table[idx].first == key)
4         return table[idx].second;
5     return 0;
6 }
```

Also very simple, we always know the value will be in position returned by `idx`. The **remove** will be of the same simplicity, as following:

```
1 void remove(string key) {
2     unsigned int idx = hashFunction(key);
3     table[idx].first = "";
4 }
```

Here we make the assumption that an empty position has an empty string. We could also carry a boolean (usually called a tombstone) to check if the position is occupied or not.

### 3.1.2 Formal definition

### 3.1.3 Linear Probing

### 3.1.4 Quadratic Probing

### 3.1.5 When to resize an array

### 3.1.6 How to delete an entry

### 3.1.7 Practical advantages of Open Addressing

## 3.2 Chaining Hashing

Chaining hashing is the implementation of a hash table using linked lists. On this implementation, each bucket of the table is a pointer to a node of a linked list, that will carry the key value pair in our case. We deal with collisions with this implementation by adding a new node to the end of the list.

This implementation is considered simpler than open addressing, usually because the way of dealing with collisions is clearer. Also it is less system dependant if we consider performance (as we saw in one of the key advantages of open addressing is that it is cache friendly). That is one of the key reasons that C++ uses chaining hashing for its default implementation of `unordered_hash` [? ].

Below we can see an implementation of chaining hashing in C++11:

### 3.2.1 Simple Chaining Hashing Algorithm

```

1 class Node {
2 public:
3     string key;
4     int value;
5     Node* nxt;
6
7     Node(string KEY = "", int VALUE = 0, Node* NXT = NULL) {
8         key = KEY;
9         value = VALUE;
10        nxt = NXT;
11    }
12
13    ~Node();
14 };
15
16 class ChainingHashTable {
17 public:

```

```
18  vector<Node*> table;
19  int m, n;
20
21  ChainingHashTable() {
22      m = 16;
23      table.resize(m);
24      n = 0;
25  }
26
27  unsigned int hashFunction(string s) {...}
28
29  void insert(string key, int value) {
30      unsigned int idx = hashFunction(key);
31      if (table[idx] == NULL) {
32          table[idx] = new Node(key, value);
33      } else {
34          Node* cur = table[idx];
35          while (cur->nxt != NULL)
36              cur = cur->nxt;
37          cur->nxt = new Node(key, value);
38      }
39
40      n++;
41      resizeIfNecessary();
42  }
43
44  int find(string key) {
45      unsigned int idx = hashFunction(key);
46      Node* cur = table[idx];
47
48      while (cur->key != key)
49          cur = cur->nxt;
50
51      if (cur->key == key)
52          return cur->value;
53      return 0;
54  }
55
56  void remove(string key) {
57      unsigned int idx = hashFunction(key);
58      // Empty bucket
59      if (table[idx] == NULL)
60          return;
61
62      // Remove head of bucket
63      if (table[idx]->key == key) {
64          Node* rem = table[idx];
65          table[idx] = table[idx]->nxt;
```

```

66         free (rem);
67         return;
68     }
69
70     // General case
71     Node* cur = table[idx];
72     while (cur->nxt != NULL && cur->nxt->key != key)
73         cur = cur->nxt;
74     if (cur->nxt != NULL) {
75         Node* rem = cur->nxt;
76         cur->nxt = rem->nxt;
77         free (rem);
78     }
79
80 }
81
82 private:
83     void resizeIfNecessary() {...};
84 };

```

The implementation has omitted the implementations of **hashFunction** because that has already been discussed and **resizeIfNecessary** because that will be discussed later. The methods **insert**, **find** and **remove**

### 3.2.2 Move to front

### 3.2.3 When to resize an array

### 3.2.4 How to delete an entry

### 3.2.5 Practical advantages of Chaining Hashing

# Chapter 4

## Applications

- Give a glance at what type of application we have, focus on 2 algorithmic
- Rabin Karp
- Hashing trees

Hash functions and hash tables have a great number of applications in computer science. During this last section I focus on applications of hash functions in algorithms, but citing superficially applications in other areas (like cryptography, data deduplication and caching).

Among the applications that I explain in this section there is a focus in two applications: Rabin-Karp string matching algorithm and hashing of a rooted tree for isomorphism checking. Rabin-karp string matching algorithm is one of the main application of a technique called rolling hashing. Hashing of rooted tree for isomorphism checking is an interesting application sometimes used in competitive programming.

To motivate the start of this thesis I will start using hash tables to solve a very simple, yet famous, problem called 3-sum.

### 4.1 3-sum problem

The problem is stated as following:

*Make a function that given an array of integer numbers and an integer  $S$ , it returns if there are any 3 different elements in this array that its sum equals  $S$ . Assume that there are no three different elements in the array that overflow a 32-bit integer when summed together.*

This is a very interesting problem that has many different solutions. To start I will show and explain to you the brute force solution:

```

1 bool threeSumWithoutHashTable(vector<int> v, int S) {
2     for (int i = 0; i < v.size(); i++) {
3         for (int j = i + 1; j < v.size(); j++) {
4             for (int k = j + 1; k < v.size(); k++) {
5                 if (v[i] + v[j] + v[k] == S) return true;
6             }
7         }
8     }
9     return false;
10 }

```

The above solution solves the problem in  $O(n^3)$  time complexity and  $O(1)$  memory complexity, being  $n$  the size of the array. It doesn't allocate any memory but checks every triple to find if one satisfies the condition. The question is, can we do better in time complexity using hash tables? The answer is yes:

```

1 bool threeSumWithHashTable(vector<int> v, int S) {
2     unordered_map<int, int> hashTable;
3     for (int i = 0; i < v.size(); i++) {
4         hashTable[v[i]]++;
5     }
6     for (int i = 0; i < v.size(); i++) {
7         for (int j = i + 1; j < v.size(); j++) {
8             hashTable[v[i]]--;
9             hashTable[v[j]]--;
10            if (hashTable.find(S - v[i] - v[j]) != hashTable.end() &&
11                hashTable[S - v[i] - v[j]] > 0) return true;
12            hashTable[v[i]]++;
13            hashTable[v[j]]++;
14        }
15    }
16    return false;
17 }

```

The above solution solves the problem in  $O(n^2)$ , time complexity (average) and  $O(n)$  memory complexity. Although the worst case scenario is  $O(n^3)$  and it uses more memory, this solution is way faster in practice for large input cases. To showcase this I made simulations with the codes shown (that can be found in bibliography), the results are:



ArraySize	Time Without Hash Table	Time with Hash Table	Increase in Performance
128	4.231ms	6.494ms	-53.4%
256	34.223ms	26.665ms	22.0%
512	267.499ms	99.130ms	62.9%
1024	1742.688ms	302.453ms	82.6%
2048	7345.126ms	683.197ms	90.6%
4096	25029.888ms	761.363ms	96.9%

As we can see for the table above, the three sum solution using hash table quickly surpasses the brute force implementation. As we will see later, hash tables (or *unordered\_map*) are not the fastest hash tables possible [? ]. That means we can have an even greater performance than what is shown now.



# Chapter 5

## Conclusions

[illegible]

<sup>1</sup>Exemplo de referência para página Web: [www.vision.ime.usp.br/~jmena/stuff/tese-exemplo](http://www.vision.ime.usp.br/~jmena/stuff/tese-exemplo)



## Appendix A

# Appendix

[illegible]

